

Database System Principle - SQL Query Language

李旭东

LEEXUDONG@NANKAI.EDU.CN
NANKAI UNIVERSITY

OBJECTIVES

- Overview of The SQL Query Language
 - Data Definition
 - Basic Query Structure
 - Additional Basic Operations
 - Set Operations
 - Null Values
 - Aggregate Functions
 - Nested Subqueries
 - Modification of the Database
- ©LXD

OVERVIEW OF THE SQL QUERY LANGUAGE

- IBM **Sequel** language developed as part of System R project at the IBM San Jose Research Laboratory
 - 1970s
 - Renamed **Structured Query Language (SQL)**
 - In 1986, the American National Standards Institute (ANSI) and the International Organization for Standardization (ISO) published an SQL standard, called SQL-86
- ©LXD

ANSI AND ISO STANDARD SQL

- SQL-86 (or SQL-87) is the ISO 9075:1987 standard of 1987
 - SQL-89 is the ISO/IEC 9075:1989 standard of 1989
 - SQL-92 is the ISO/IEC 9075:1992 standard of 1992
 - SQL:1999 is the ISO/IEC 9075:1999 standard of 1999
 - SQL:2003 is the ISO/IEC 9075:2003 standard of 2003*
 - SQL:2006 is the ISO/IEC 9075:2006 standard of 2006
 - SQL:2008 is the ISO/IEC 9075:2008 standard of 2008
 - SQL:2011 is the ISO/IEC 9075:2011 standard of 2011
- ©LXD

SQL2003

- 1) ISO/IEC 9075—1: Framework(SQL/Framework)
 - 2) ISO/IEC 9075—2: Foundation(SQL/Foundation)
 - 3) ISO/IEC 9075—3: Call Level Interface(SQL/CLI)
 - 4) ISO/IEC 9075—4: Persistent Stored Modules(SQL/PSM)
 - 5) ISO/IEC 9075—9: Management of External Data(SQL/MED)
 - 6) ISO/IEC 9075—10: Object Language Bindings(SQL/OLB)
 - 7) ISO/IEC 9075—11: Information and Definition Schemas(SQL/Schemata)
 - 8) ISO/IEC 9075—13: Java Routines and Types Using the Java Programming Language(SQL/JRT)
 - 9) ISO/IEC 9075—14: XML Related Specifications(SQL/XML)
- ©LXD

OVERVIEW OF THE SQL QUERY LANGUAGE

- Structured Query Language (SQL)
 - Data-Definition Language, DDL
 - Data-Manipulation Language, DML
 - Integrity完整性
 - View Definition视图定义
 - Transaction Control事务控制
 - Embedded SQL and dynamic SQL嵌入式和动态SQL
 - Authorization授权
 - Commercial systems offer most, if not all, **SQL-92** features, plus varying feature sets from later standards and special proprietary features.
- ©LXD

OBJECTIVES

- Overview of The SQL Query Language
- **Data Definition**
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

SQL DATA DEFINITION

- **SQL DDL**
 - The schema for each relation
 - The types of values associated with each attribute
 - The integrity constraints
 - The set of indices to be maintained for each relation
 - The security and authorization information for each relation
 - The physical storage structure of each relation on disk

©LXD

DOMAIN TYPES IN SQL

- **char(n)**. Fixed length character string, with user-specified length n.
- **varchar(n)**. Variable length character strings, with user-specified maximum length n.

©LXD


DOMAIN TYPES IN SQL

- **int**. Integer (a finite subset of the integers that is machine-dependent).
- **smallint**. Small integer (a machine-dependent subset of the integer domain type).
- **numeric(p,d)**. Fixed point number, with user-specified precision of p digits, with d digits to the right of decimal point.
 - (ex., numeric(3,1), allows 44.5, not 444.5 or 0.32)
- ...

©LXD

DOMAIN TYPES IN SQL

- **real, double** precision. Floating point 浮点 and double-precision 双精度 floating point numbers, with *machine-dependent precision*
 - 1 bit (sign bit), 11 bits (exponential bit), 52 bits (tail digit bit), $2^{52} = 4503599627370496$
 - 16bits. double的精度为15~16位



- **float(n)**. 单精度 Floating point number, with user-specified precision of at least n digits
 - 1 bit (sign bit) 8 bits (exponential 指数 bit) 23 bits (tail digit bit), $2^{23} = 8388608$, 7bits, float的精度为6~7位有效数字

©LXD

DATA TYPE IN MYSQL

- **Numeric Data Types**
 - Integer Types (Exact Value)
 - INTEGER, INT, SMALLINT, TINYINT, MEDIUMINT, **BIGINT**
 - Fixed-Point Types (Exact Value)
 - DECIMAL 十进制的, NUMERIC
 - Floating-Point Types (Approximate Value)
 - FLOAT, DOUBLE
 - Bit-Value Type - BITS

©LXD

DATA TYPE IN MYSQL

- Date and Time Data Types
 - The DATE Type: '1000-01-01' to '9999-12-31'
 - The DATETIME Type: 'YYYY-MM-DD hh:mm:ss', 'YYYY-MM-DD hh:mm:ss[.fraction]'
 - '1000-01-01 00:00:00.000000' to '9999-12-31 23:59:59.999999'
 - The TIMESTAMP Type: 'YYYY-MM-DD hh:mm:ss[.fraction]'
 - '1970-01-01 00:00:01.000000' to '2038-01-19 03:14:07.499999'
 - The TIME Type: 'hh:mm:ss': '-838:59:59' to '838:59:59'
 - The YEAR Type: YYYY format, with a range of 1901 to 2155, and 0000
- Automatic Initialization and Updating for TIMESTAMP and DATETIME

DATA TYPE IN MYSQL

- String Data Types
 - The CHAR and VARCHAR Types
 - The BINARY and VARBINARY Types
 - The BLOB and TEXT Types
 - The ENUM Type
 - ENUM('x-small', 'small', 'medium', 'large', 'x-large')
 - The SET Type

BASIC SCHEMA DEFINITION

- Create table
- Insert into
- Delete from
- Alter table
- Drop table

CREATE TABLE CONSTRUCT

- An SQL relation is defined using the **create table** command:

```
create table r (A1 D1,  
               A2 D2, ...,  
               An Dn,  
               (integrity-constraint1 完整性约束),  
               ...,  
               (integrity-constrainti))
```

- r is the name of the relation
- each A_i is an attribute name in the schema of relation r
- D_i is the data type of values in the domain of attribute A_i

CREATE TABLE CONSTRUCT

- the **create table** command
- Example:

```
create table instructor (  
    ID        char(5),  
    name      varchar(20),  
    dept_name varchar(20),  
    salary    numeric(8,2));
```

INTEGRITY CONSTRAINTS IN CREATE TABLE

- not null
- primary key ($A_1, \dots, A_n$)
- foreign key ($A_m, \dots, A_n$) references r

Example:

```
create table instructor (  
    ID        char(5),  
    name      varchar(20) not null,  
    dept_name varchar(20),  
    salary    numeric(8,2),  
    primary key (ID),  
    foreign key (dept_name) references department);
```

primary key declaration on an attribute automatically ensures **not null**

CASES: CREATE TABLE

```
create table department
(dept_name varchar (20),
building varchar (15),
budget numeric (12,2),
primary key (dept_name);

create table course
(course_id varchar (7),
title varchar (50),
dept_name varchar (20),
credits numeric (2,0),
primary key (course_id),
foreign key (dept_name) references department(dept_name));
```

©LXD

CASES: CREATE TABLE

```
create table section
(course_id varchar (8),
sec_id varchar (8),
semester varchar (6),
year numeric (4,0),
building varchar (15),
room_number varchar (7),
time_slot_id varchar (4),
primary key (course_id, sec_id, semester, year),
foreign key (course_id) references course);
```

©LXD

CASES: CREATE TABLE

```
create table teaches
( ID varchar (5),
course_id varchar (8),
sec_id varchar (8),
semester varchar (6),
year numeric (4,0),
primary key ( ID , course_id, sec_id, semester, year),
foreign key (course_id, sec_id, semester, year) references section,
foreign key ( ID ) references instructor);
```

©LXD

UPDATES TO TABLES

- **Insert**
 - **insert into** *instructor* values ('10211', 'Smith', 'Biology', 66000);
- **Delete**
 - Remove all tuples from the *student* relation
 - **delete from** *student*
- **Drop Table**
 - **drop table** *r*

©LXD

UPDATES TO TABLES

- **Alter**
 - **alter table** *r* **add** *A* *D*
 - where *A* is the name of the attribute to be added to relation *r* and *D* is the domain of *A*.
 - All exiting tuples in the relation are assigned *null* as the value for the new attribute.
 - **alter table** *r* **drop** *A*
 - where *A* is the name of an attribute of relation *r*
 - Dropping of attributes not supported by many databases.

©LXD

OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- **Basic Query Structure**
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

BASIC STRUCTURE OF SQL QUERIES

- A typical SQL query has the form:

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

A_i represents an attribute

- R_i represents a relation
 - P is a predicate.
- The result of an SQL query is a **relation**.

©LXD

QUERIES ON A SINGLE RELATION

- The **select** clause lists the attributes desired in the result of a query
 - corresponds to the projection operation of the relational algebra

- Example: find the names of all instructors:

```
select name
from instructor
```

- NOTE: SQL names are case insensitive
 - E.g., *Name* ≡ *NAME* ≡ *name*
 - Some people use upper case wherever we use bold font.

name
Srinivasan
Wu
Mozart
Einstein
El Said
Gold
Katz
Califieri
Singh
Crick
Brandt
Kim

©LXD

QUERIES ON A SINGLE RELATION

- DEDUPLICATION 去重

- SQL allows duplicates in relations as well as in query results.
- To force the elimination of duplicates, insert the keyword **distinct** after select.
- Find the department names of all instructors, and remove duplicates

```
select distinct dept_name
from instructor
```

- The keyword **all** specifies that duplicates should not be removed.

```
select all dept_name
from instructor
```

©LXD

QUERIES ON A SINGLE RELATION

- MORE

- An **asterisk (*)** in the select clause denotes "all attributes"

```
select *
from instructor
```

- An attribute can be a **literal** with no **from** clause

```
select '437'
```

- Results is a table with one column and a single row with value "437"
- Can give the column a name using:

```
select '437' as FOO
```

©LXD

QUERIES ON A SINGLE RELATION

- MORE

- An attribute can be a literal with **from** clause

```
select 'A'
from instructor
```

- Result is a table with **one column** and **N rows** (number of tuples in the *instructors* table), each row with value "A"

©LXD

QUERIES ON A SINGLE RELATION

- MORE

- The **select** clause can contain arithmetic expressions involving the operation, +, -, *, and /, and operating on constants or attributes of tuples.

- The query:

```
select ID, name, salary/12
from instructor
```

would return a relation that is the same as the *instructor* relation, except that the value of the attribute *salary* is divided by 12.

- Can rename "*salary/12*" using the **as** clause:

```
select ID, name, salary/12 as monthly_salary
```

©LXD

THE WHERE CLAUSE

- The **where** clause specifies conditions that the result must satisfy
 - Corresponds to the selection predicate of the relational algebra.
- To find all instructors in Comp. Sci. dept

```
select name
from instructor
where dept_name = 'Comp. Sci.'
```

©LXD

THE WHERE CLAUSE

- Comparison results can be combined using the logical connectives **and**, **or**, and **not**
 - To find all instructors in Comp. Sci. dept with salary > 80000


```
select name
from instructor
where dept_name = 'Comp. Sci.' and salary > 80000
```
- Comparisons can be applied to results of arithmetic expressions.

©LXD

QUERIES ON MULTIPLE RELATIONS

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P;
```

- Each **A_i** represents an attribute, and each **r_i** a relation.
- P is a predicate
 - If the where clause is omitted, the predicate P is true.

©LXD

QUERIES ON MULTIPLE RELATIONS

- The **from** clause lists the relations involved in the query
 - Corresponds to the Cartesian product operation of the relational algebra.
- Find the Cartesian product *instructor* X *teaches*

```
select *
from instructor, teaches
```

 - generates every possible instructor – teaches pair, with all attributes from both relations.
 - For common attributes (e.g., *ID*), the attributes in the resulting table are renamed using the relation name (e.g., *instructor.ID*)

©LXD

CARTESIAN PRODUCT 笛卡尔积

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

inst.ID	name	dept_name	salary	teaches.ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	10101	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	10101	CS-347	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	12121	FIN-201	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	15151	MU-199	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...
12121	Wu	Finance	90000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
12121	Wu	Finance	90000	10101	CS-347	1	Fall	2009
12121	Wu	Finance	90000	12121	FIN-201	1	Spring	2010
12121	Wu	Finance	90000	15151	MU-199	1	Spring	2010
12121	Wu	Finance	90000	22222	PHY-101	1	Fall	2009
...	...	...	...	...	...	...	...	...

©LXD

CASES: CARTESIAN PRODUCT

- Find the names of all instructors who have taught some course and the course_id
 - ```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID
```
- Find the names of all instructors in the Art department who have taught some course and the course\_id
  - ```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID and instructor.dept_name = 'Art'
```

©LXD

THE NATURAL JOIN

```
select A1, A2, ..., An
from r1 natural join r2 natural join ... natural join rm
where P;
```

More generally, a **from** clause can be of the form

```
from E1, E2, ..., En
```

- where each E_i can be a single relation or an expression involving natural joins.

CASES: THE NATURAL JOIN

The natural join of the instructor relation with the teaches relation

ID	name	dept_name	salary
10101	Srinivasan	Comp. Sci.	65000
12121	Wu	Finance	90000
15151	Mozart	Music	40000
22222	Einstein	Physics	95000
32343	El Said	History	60000

ID	course_id	sec_id	semester	year
10101	CS-101	1	Fall	2009
10101	CS-315	1	Spring	2010
10101	CS-347	1	Fall	2009
12121	FIN-201	1	Spring	2010
15151	MU-199	1	Spring	2010
22222	PHY-101	1	Fall	2009

ID	name	dept_name	salary	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	CS-101	1	Fall	2009
10101	Srinivasan	Comp. Sci.	65000	CS-315	1	Spring	2010
10101	Srinivasan	Comp. Sci.	65000	CS-347	1	Fall	2009
12121	Wu	Finance	90000	FIN-201	1	Spring	2010
15151	Mozart	Music	40000	MU-199	1	Spring	2010
22222	Einstein	Physics	95000	PHY-101	1	Fall	2009
32343	El Said	History	60000	HIS-351	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-101	1	Spring	2010
45565	Katz	Comp. Sci.	75000	CS-319	1	Spring	2010
76766	Crick	Biology	72000	BIO-101	1	Summer	2009
76766	Crick	Biology	72000	BIO-301	1	Summer	2010
83821	Brandt	Comp. Sci.	92000	CS-190	1	Spring	2009
83821	Brandt	Comp. Sci.	92000	CS-190	2	Spring	2009
83821	Brandt	Comp. Sci.	92000	CS-319	2	Spring	2010
98345	Kim	Elec. Eng.	80000	EE-181	1	Spring	2009

CASES: THE NATURAL JOIN

• Way 1

```
select name, course_id
from instructor, teaches
where instructor.ID = teaches.ID;
```

• Way 2

```
select name, course_id
from instructor natural join teaches;
```

CASES 2: THE NATURAL JOIN

```
select name, title
```

```
1. from instructor natural join teaches, course
   where teaches.course_id = course.course_id;
```

```
2. select name, title
   from instructor natural join teaches natural join course;
```

```
3. select name, title
   from (instructor natural join teaches) join course using (course_id);
```

CASES 2: THE NATURAL JOIN

-> SCHEMA OF THE UNIVERSITY DATABASE

```
classroom(building, room_number, capacity)
department(dept_name, building, budget)
course(course_id, title, dept_name, credits)
instructor(ID, name, dept_name, salary)
section(course_id, sec_id, semester, year, building, room_number, time_slot_id)
teaches(ID, course_id, sec_id, semester, year)
student(ID, name, dept_name, tot_cred)
takes(ID, course_id, sec_id, semester, year, grade)
advisor(s_ID, i_ID)
time_slot(time_slot_id, day, start_time, end_time)
prereq(course_id, prereq_id)
```

OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- **Additional Basic Operations**
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

THE RENAME OPERATION

- The SQL allows renaming relations and attributes using the **as** clause:

old-name as new-name

- Find the names of all instructors who have a higher salary than some instructor in 'Comp. Sci'.

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept_name = 'Comp. Sci.'
```

- Keyword **as** is optional and may be omitted

instructor as T $\equiv$ *instructor T*

STRING OPERATIONS

- SQL includes a string-matching operator for comparisons on character strings. The operator **like** uses patterns that are described using two special characters:

- percent (%). The % character matches any substring.
- underscore (_). The _ character matches any character.

STRING OPERATIONS

- Find the names of all instructors whose name includes the substring "dar".

```
select name
from instructor
where name like '%dar%'
```

- Match the string "100%"

```
like '100 \%' escape '\'
```

in that above we use backslash (\) as the escape character.

STRING OPERATIONS

- Patterns are case sensitive.
- Pattern matching examples:
 - 'Intro%' matches any string beginning with "Intro".
 - '%Comp%' matches any string containing "Comp" as a substring.
 - '___' matches any string of exactly three characters.
 - '___%' matches any string of at least three characters.

STRING OPERATIONS

- SQL supports a variety of string operations such as

- concatenation (using "||")
- converting from upper to lower case (and vice versa)
 - lower(), upper()
- finding string length, extracting substrings, etc.
 - LENGTH()
 - CHAR_LENGTH()
 - Substr(str,pos[,length])

ORDERING THE DISPLAY OF TUPLES

- List in alphabetic order the names of all instructors

```
select distinct name
from instructor
order by name
```

- We may specify **desc** for descending order or **asc** for ascending order, for each attribute; ascending order is the default.

- Example: **order by name desc**
- Can sort on multiple attributes
 - Example: **order by dept_name, name**

WHERE CLAUSE PREDICATES

- SQL includes a **between** comparison operator
- Example: Find the names of all instructors with salary between \$90,000 and \$100,000 (that is, $\geq \$90,000$ and $\leq \$100,000$)
 - ```
select name
from instructor
where salary between 90000 and 100000
```

©LXD

## WHERE CLAUSE PREDICATES

- Way 1:
  - ```
select name, course_id
from instructor, teaches
where instructor.ID= teaches.ID and dept_name = 'Biology';
```
- Way 2:
 - ```
select name, course_id
from instructor, teaches
where (instructor.ID, dept_name) = (teaches.ID, 'Biology');
```

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- **Set Operations**
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## SET OPERATIONS

- union
- intersect
- except

©LXD

## SET OPERATIONS

### - UNION

- Find courses that ran in Fall 2009 or in Spring 2010

```
(select course_id from section where sem = 'Fall' and
year = 2009)
union
(select course_id from section where sem = 'Spring' and
year = 2010)
```

©LXD

## SET OPERATIONS

### - INTERSECT

- Find courses that ran in Fall 2009 and in Spring 2010

```
(select course_id from section where sem = 'Fall' and
year = 2009)
intersect
(select course_id from section where sem = 'Spring'
and year = 2010)
```

©LXD

## SET OPERATIONS

### - EXCEPT

- Find courses that ran in Fall 2009 but not in Spring 2010

```
(select course_id from section where sem = 'Fall' and
year = 2009)
```

**except**

```
(select course_id from section where sem = 'Spring'
and year = 2010)
```

©LXD

## SET OPERATIONS

- Set operations **union**, **intersect**, and **except**
  - Each of the above operations automatically eliminates duplicates
- To retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.
- Suppose a tuple occurs  $m$  times in  $r$  and  $n$  times in  $s$ , then, it occurs:
  - $m + n$  times in  **$r$  union all  $s$**
  - $\min(m, n)$  times in  **$r$  intersect all  $s$**
  - $\max(0, m - n)$  times in  **$r$  except all  $s$**

©LXD

读范孙旧诗文兼怀古

AL 20220310

风云百变世消磨，衰骨经年叹几何？  
运去芳华怀旧梦，时来天地寄新歌。  
津门纵笔添愁绪，南浦移舟起浪波。  
物换人非寻所住，相识唯有半池荷。

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- **Null Values**
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD

## NULL VALUES

- It is possible for tuples to have a null value, denoted by **null**, for some of their attributes
- **null** signifies an unknown value or that a value does not exist.
- The result of any arithmetic expression involving **null** is **null**
  - Example:  $5 + \text{null}$  returns null
- The predicate **is null** can be used to check for null values.
  - Example: Find all instructors whose salary is null.

```
select name
from instructor
where salary is null
```

©LXD

## NULL VALUES AND THREE VALUED LOGIC

- Three logic values – **true**, **false**, **unknown**
- Any comparison with **null** returns **unknown**
  - Example:  $5 < \text{null}$  or  $\text{null} <> \text{null}$  or  $\text{null} = \text{null}$

©LXD

## NULL VALUES AND THREE VALUED LOGIC

- Three-valued logic using the value *unknown*:
  - OR: (*unknown or true*) = *true*,  
(*unknown or false*) = *unknown*  
(*unknown or unknown*) = *unknown*
  - AND: (*true and unknown*) = *unknown*,  
(*false and unknown*) = *false*,  
(*unknown and unknown*) = *unknown*
  - NOT: (*not unknown*) = *unknown*
  - “*P is unknown*” evaluates to *true* if predicate *P* evaluates to *unknown*
- Result of **where** clause predicate is treated as *false* if it evaluates to *unknown*

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- **Aggregate Functions**
- Nested Subqueries
- Modification of the Database

©LXD

## AGGREGATE FUNCTIONS 聚集函数

- These functions operate on the multiset of values of a column of a relation, and return a value

**avg**: average value  
**min**: minimum value  
**max**: maximum value  
**sum**: sum of values  
**count**: number of values

©LXD

## AGGREGATE FUNCTIONS

- Find the average salary of instructors in the Computer Science department
  - **select avg (salary)**  
**from instructor**  
**where dept\_name= 'Comp. Sci.';**
- Find the total number of instructors who teach a course in the Spring 2010 semester
  - **select count (distinct ID)**  
**from teaches**  
**where semester = 'Spring' and year = 2010;**

©LXD

## AGGREGATE FUNCTIONS – GROUP BY

- Find the average salary of instructors in each department
  - **select dept\_name, avg (salary) as avg\_salary**  
**from instructor**  
**group by dept\_name;**

| ID    | name       | dept_name  | salary |
|-------|------------|------------|--------|
| 76766 | Crick      | Biology    | 72000  |
| 45565 | Katz       | Comp. Sci. | 75000  |
| 10101 | Srinivasan | Comp. Sci. | 65000  |
| 83821 | Brandt     | Comp. Sci. | 92000  |
| 98345 | Kim        | Elec. Eng. | 80000  |
| 12121 | Wu         | Finance    | 90000  |
| 76543 | Singh      | Finance    | 80000  |
| 32343 | El Said    | History    | 60000  |
| 58583 | Califieri  | History    | 62000  |
| 15151 | Mozart     | Music      | 40000  |
| 33456 | Gold       | Physics    | 87000  |
| 22222 | Einstein   | Physics    | 95000  |

©LXD

| dept_name  | salary |
|------------|--------|
| Biology    | 72000  |
| Comp. Sci. | 77333  |
| Elec. Eng. | 80000  |
| Finance    | 85000  |
| History    | 61000  |
| Music      | 40000  |
| Physics    | 91000  |

## AGGREGATE FUNCTIONS – GROUP BY

- Attributes in **select** clause outside of aggregate functions must appear in **group by** list
  - /\* **erroneous** query \*/  
**select dept\_name, ID, avg (salary)**  
**from instructor**  
**group by dept\_name;**

©LXD



## AGGREGATE FUNCTIONS – HAVING CLAUSE

- Find the names and average salaries of all departments whose average salary is greater than 42000

```
select dept_name, avg (salary)
from instructor
group by dept_name
having avg (salary) > 42000;
```

Note: predicates in the **having** clause are applied after the formation of groups whereas predicates in the **where** clause are applied before forming groups

## AGGREGATE FUNCTIONS

- Total all salaries

```
select sum (salary)
from instructor
```

- Above statement ignores null amounts
- Result is *null* if there is no non-null amount

## AGGREGATE FUNCTIONS

- All aggregate operations except **count(\*)** ignore tuples with null values on the aggregated attributes
- What if collection has only null values?
  - count returns 0
  - all other aggregates return null

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries**
- Modification of the Database

## NESTED SUBQUERIES 嵌套子查询

- SQL provides a mechanism for nesting subqueries.
- A subquery is a select-from-where expression that is nested within another query.

```
select distinct course_id
from section
where semester = 'Fall' and year= 2009 and
course_id in (select course_id
from section
where semester = 'Spring' and year= 2010);
```

## NESTED SUBQUERIES 嵌套子查询

- The nesting can be done in the following SQL query

```
select A1, A2, ..., An
from r1, r2, ..., rm
where P
```

as follows:

- A<sub>i</sub> can be replaced by a subquery that generates a single value.
- r<sub>i</sub> can be replaced by any valid subquery
- P can be replaced with an expression of the form:  
B <operation> (subquery)

Where B is an attribute and <operation> to be defined later.

## SUBQUERIES IN THE WHERE CLAUSE

- A common use of subqueries is to perform tests:
  - For set membership
  - For set comparisons
  - For set cardinality 基数

## NESTED SUBQUERIES

### - SET MEMBERSHIP 集合成员测试

- in, not in

- Find courses offered in Fall 2009 but not in Spring 2010

```
select distinct course_id
from section
where semester = 'Fall' and year = 2009 and
 course_id not in (select course_id
 from section
 where semester = 'Spring' and year = 2010);
```

## NESTED SUBQUERIES

### - SET MEMBERSHIP

- Find the total number of (distinct) students who have taken course sections taught by the instructor with ID 10101

```
select count (distinct ID)
from takes
where (course_id, sec_id, semester, year) in
 (select course_id, sec_id, semester, year
 from teaches
 where teaches.ID = 10101);
```

## NESTED SUBQUERIES

### - SET COMPARISON 集合比较

- Find names of instructors with salary greater than that of some (at least one) instructor in the Biology department.

```
select distinct T.name
from instructor as T, instructor as S
where T.salary > S.salary and S.dept_name = 'Biology';
```

- Same query using > some clause

```
select name
from instructor
where salary > some (select salary
 from instructor
 where dept_name = 'Biology');
```

## NESTED SUBQUERIES

### - SET COMPARISON 集合比较

- > some

- < some

- <= some

- >= some

- <> some

- (not same to not in)

- =some (same to in)

- $F \langle \text{comp} \rangle \text{some } r \Leftrightarrow \exists t \in r \text{ such that } (F \langle \text{comp} \rangle t)$

Where <comp> can be: <, ≤, >, =, ≠

|           |                                                                                       |   |   |                      |        |                                        |
|-----------|---------------------------------------------------------------------------------------|---|---|----------------------|--------|----------------------------------------|
| (5 < some | <table border="1"><tr><td>0</td></tr><tr><td>5</td></tr><tr><td>6</td></tr></table> ) | 0 | 5 | 6                    | = true | (read: 5 < some tuple in the relation) |
| 0         |                                                                                       |   |   |                      |        |                                        |
| 5         |                                                                                       |   |   |                      |        |                                        |
| 6         |                                                                                       |   |   |                      |        |                                        |
| (5 < some | <table border="1"><tr><td>0</td></tr><tr><td>5</td></tr></table> )                    | 0 | 5 | = false              |        |                                        |
| 0         |                                                                                       |   |   |                      |        |                                        |
| 5         |                                                                                       |   |   |                      |        |                                        |
| (5 = some | <table border="1"><tr><td>0</td></tr><tr><td>5</td></tr></table> )                    | 0 | 5 | = true               |        |                                        |
| 0         |                                                                                       |   |   |                      |        |                                        |
| 5         |                                                                                       |   |   |                      |        |                                        |
| (5 ≠ some | <table border="1"><tr><td>0</td></tr><tr><td>5</td></tr></table> )                    | 0 | 5 | = true (since 0 ≠ 5) |        |                                        |
| 0         |                                                                                       |   |   |                      |        |                                        |
| 5         |                                                                                       |   |   |                      |        |                                        |

(= some) = in  
However, (≠ some) ≠ not in

## NESTED SUBQUERIES

### - SET COMPARISON 集合比较

- > all

- < all

- <= all

- >= all

- =all (not same to in)

- <> all (same to not in)

- $F \langle \text{comp} \rangle \text{all } r \Leftrightarrow \forall t \in r (F \langle \text{comp} \rangle t)$

|          |                                                                                       |   |    |                                |         |
|----------|---------------------------------------------------------------------------------------|---|----|--------------------------------|---------|
| (5 < all | <table border="1"><tr><td>0</td></tr><tr><td>5</td></tr><tr><td>6</td></tr></table> ) | 0 | 5  | 6                              | = false |
| 0        |                                                                                       |   |    |                                |         |
| 5        |                                                                                       |   |    |                                |         |
| 6        |                                                                                       |   |    |                                |         |
| (5 < all | <table border="1"><tr><td>6</td></tr><tr><td>10</td></tr></table> )                   | 6 | 10 | = true                         |         |
| 6        |                                                                                       |   |    |                                |         |
| 10       |                                                                                       |   |    |                                |         |
| (5 = all | <table border="1"><tr><td>4</td></tr><tr><td>5</td></tr></table> )                    | 4 | 5  | = false                        |         |
| 4        |                                                                                       |   |    |                                |         |
| 5        |                                                                                       |   |    |                                |         |
| (5 ≠ all | <table border="1"><tr><td>4</td></tr><tr><td>6</td></tr></table> )                    | 4 | 6  | = true (since 5 ≠ 4 and 5 ≠ 6) |         |
| 4        |                                                                                       |   |    |                                |         |
| 6        |                                                                                       |   |    |                                |         |

(≠ all) = not in  
However, (= all) ≠ in

## NESTED SUBQUERIES

### - SET COMPARISON 集合比较

?What meaning?

```
select dept_name
from instructor
group by dept_name
having avg (salary) >= all (select avg (salary)
 from instructor
 group by dept_name);
```

©LXD

## TEST FOR EMPTY RELATIONS

- exists , not exists
- Yet another way of specifying the query "Find all courses taught in both the Fall 2009 semester and in the Spring 2010 semester"

```
select course_id
from section as S
where semester = 'Fall' and year = 2009 and
 exists (select *
 from section as T
 where semester = 'Spring' and year = 2010 and S.course_id = T.course_id);
```

- **Correlation name** – variable S in the outer query
- **Correlated subquery** – the inner query

©LXD

## TEST FOR EMPTY RELATIONS

- Find all students who have taken all courses offered in the Biology department.

```
select distinct S.ID, S.name
from student as S
where not exists ((select course_id from course
 where dept_name = 'Biology')
 except
 (select T.course_id from takes as T
 where S.ID = T.ID));
```

- First nested query lists all courses offered in Biology
- Second nested query lists all courses a particular student took
- Note that  $X - Y = \emptyset \Leftrightarrow X \subseteq Y$
- Note: Cannot write this query using = all and its variants

©LXD

## TEST FOR ABSENCE OF DUPLICATE TUPLES

- The **unique** construct tests whether a subquery has any duplicate tuples in its result.
- The **unique** construct evaluates to "true" if a given subquery contains no duplicates .
- Find all courses that were offered at most once in 2009

```
select T.course_id
from course as T
where unique (select R.course_id
 from section as R
 where T.course_id = R.course_id
 and R.year = 2009);
```

©LXD

## SUBQUERIES IN THE FROM CLAUSE

- SQL allows a subquery expression to be used in the **from** clause
- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."

```
select dept_name, avg_salary
from (select dept_name, avg (salary) as avg_salary
 from instructor
 group by dept_name)
where avg_salary > 42000;
```

- Note that we do not need to use the **having** clause

©LXD

## SUBQUERIES IN THE FROM CLAUSE

- Find the average instructors' salaries of those departments where the average salary is greater than \$42,000."
- Another way to write above query

```
select dept_name, avg_salary
from (select dept_name, avg (salary)
 from instructor
 group by dept_name) as dept_avg (dept_name, avg_salary)
where avg_salary > 42000;
```

©LXD



## WITH CLAUSE

- The **with** clause provides a way of defining a temporary relation whose definition is available only to the query in which the **with** clause occurs
  - Introduced in SQL:1999
- Find all departments with the maximum budget

```
with max_budget (value) as
(select max(budget)
 from department)
select department.name
from department, max_budget
where department.budget = max_budget.value;
```

©LXD

## COMPLEX QUERIES USING WITH CLAUSE

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total (dept_name, value) as
(select dept_name, sum(salary)
 from instructor
 group by dept_name),
dept_total_avg(value) as
(select avg(value)
 from dept_total)
select dept_name
from dept_total, dept_total_avg
where dept_total.value > dept_total_avg.value;
```

©LXD

## SCALAR SUBQUERIES 标量子查询

- Scalar subqueries**
  - SQL allows subqueries to occur wherever an expression **returning a value** is permitted, provided the subquery returns only one tuple containing a single attribute
  - List all departments along with the number of instructors in each department

```
select dept_name,
(select count(*)
 from instructor where department.dept_name = instructor.dept_name)
as num_instructors
from department;
```
  - Runtime error if subquery returns more than one result tuple

©LXD

## OBJECTIVES

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database**

©LXD

## MODIFICATION OF THE DATABASE

- Deletion of tuples from a given relation.
- Insertion of new tuples into a given relation
- Updating of values in some tuples in a given relation

©LXD

## DELETION 删除

- Delete all instructors

```
delete from instructor
```
- Delete all instructors from the Finance department

```
delete from instructor
where dept_name='Finance';
```
- Delete all tuples in the *instructor* relation for those instructors associated with a department located in the Watson building.

```
delete from instructor
where dept_name in (select dept_name
 from department
 where building = "Watson");
```

©LXD

## DELETION

- Delete all instructors whose salary is less than the average salary of instructors

```
delete from instructor
where salary < (select avg (salary)
 from instructor);
```

- Problem: as we delete tuples from deposit, the average salary changes
- Solution used in SQL:
  1. First, compute **avg** (salary) and find all tuples to delete
  2. Next, delete all tuples found above (without recomputing **avg** or retesting the tuples)

©LXD

## INSERTION插入

- Add a new tuple to *course*

```
insert into course
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- or equivalently

```
insert into course (course_id, title, dept_name, credits)
values ('CS-437', 'Database Systems', 'Comp. Sci.', 4);
```

- Add a new tuple to *student* with *tot\_creds* set to null

```
insert into student
values ('3003', 'Green', 'Finance', null);
```

©LXD

## INSERTION

- Add all instructors to the *student* relation with *tot\_creds* set to 0

```
insert into student
select ID, name, dept_name, 0
from instructor
```

- The **select from where** statement is evaluated fully before any of its results are inserted into the relation.

Otherwise queries like

```
insert into table1 select * from table1
```

would cause problem

©LXD

## UPDATES更新

- Increase salaries of instructors whose salary is over \$100,000 by 3%, and all others by a 5%

- Write two **update** statements:

```
update instructor
set salary = salary * 1.03
where salary > 100000;
```

```
update instructor
set salary = salary * 1.05
where salary <= 100000;
```

- **The order is important**
- Can be done better using the **case** statement (next slide)

©LXD

## UPDATE

### - CASE STATEMENT FOR CONDITIONAL UPDATES

- Same query as before but with **case** statement

```
update instructor
set salary = case
 when salary <= 100000 then salary * 1.05
 else salary * 1.03
 end
```

©LXD

## UPDATE

### - CASE STATEMENT FOR CONDITIONAL UPDATES

- The general form of the case statement is as follows

```
case
 when pred1 then result1
 when pred2 then result2
 ...
 when predn then resultn
 else result0
end
```

©LXD

## UPDATES WITH SCALAR SUBQUERIES

- Recompute and update tot\_creds value for all students

```
update student S
set tot_cred = (select sum(credits)
 from takes, course
 where takes.course_id = course.course_id and
 S.ID= takes.ID and takes.grade <> 'F' and
 takes.grade is not null);
```

- !Sets tot\_creds to null for students who have not taken any course

©LXD

## UPDATES WITH SCALAR SUBQUERIES

- Recompute and update tot\_creds value for all students, but sets tot\_creds to 0 for students who have not taken any course

```
update student S
set tot_cred = (select
 case
 when sum(credits) is not null then sum(credits)
 else 0
 end
 from takes, course
 where takes.course_id = course.course_id and
 S.ID= takes.ID and takes.grade <> 'F' and takes.grade is not null);
```

©LXD

```
CREATE TABLE shirts (
 name VARCHAR(40),
 size ENUM('x-small', 'small', 'medium', 'large', 'x-large')
);
```

```
INSERT INTO shirts (name, size) VALUES ('dress shirt','large'), ('t-shirt','medium'),
('polo shirt','small');
```

```
SELECT name, size FROM shirts WHERE size = 'medium';
```

```
+-----+-----+
| name | size |
+-----+-----+
| t-shirt | medium |
+-----+-----+
```

```
UPDATE shirts SET size = 'small' WHERE size = 'large';
```

```
COMMIT;
©LXD
```

## CASE

## SUMMARY

- Overview of The SQL Query Language
- Data Definition
- Basic Query Structure
- Additional Basic Operations
- Set Operations
- Null Values
- Aggregate Functions
- Nested Subqueries
- Modification of the Database

©LXD